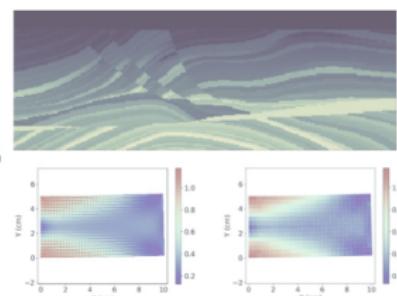
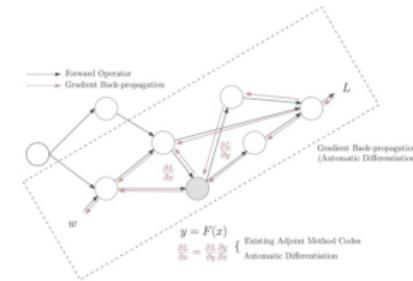
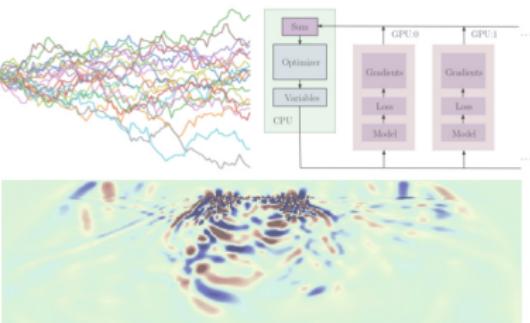


Machine Learning for Computational Engineering

Kailai Xu
Stanford University

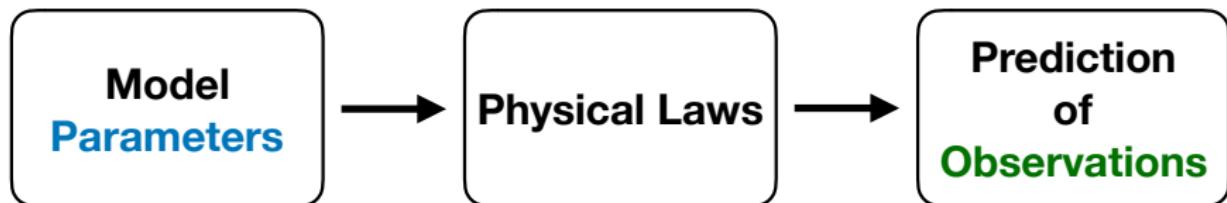


Outline

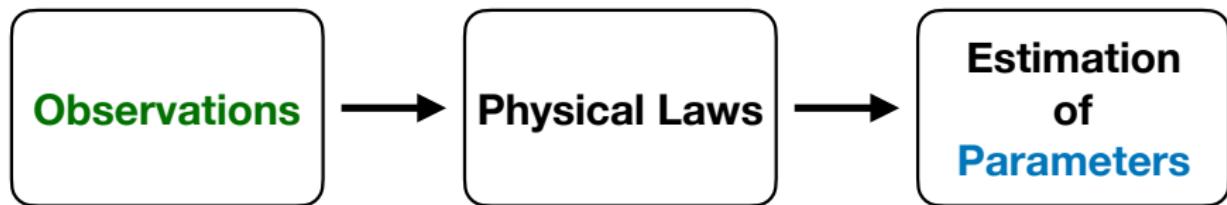
- 1 Inverse Modeling
- 2 Software Implementation
- 3 First Order Physics Constrained Learning
- 4 Second Order Physics Constrained Learning
- 5 Generative Neural Networks for Stochastic Inverse Problems
- 6 Conclusion

Inverse Modeling

Forward Problem



Inverse Problem



Inverse Modeling

We can formulate inverse modeling as a PDE-constrained optimization problem

$$\min_{\theta} L_h(u_h) \quad \text{s.t. } F_h(\theta, u_h) = 0$$

- The **loss function** L_h measures the discrepancy between the prediction u_h and the observation u_{obs} , e.g., $L_h(u_h) = \|u_h - u_{\text{obs}}\|_2^2$.
- θ is the **model parameter** to be calibrated.
- The **physics constraints** $F_h(\theta, u_h) = 0$ are described by a system of partial differential equations or differential algebraic equations (DAEs); e.g.,

$$F_h(\theta, u_h) = A(\theta)u_h - f_h = 0$$

Function Inverse Problem

$$\min_{\mathbf{f}} L_h(u_h) \quad \text{s.t. } F_h(\mathbf{f}, u_h) = 0$$

What if the unknown is a **function** instead of a set of parameters?

- Koopman operator in dynamical systems.
- Constitutive relations in solid mechanics.
- Turbulent closure relations in fluid mechanics.
- ...

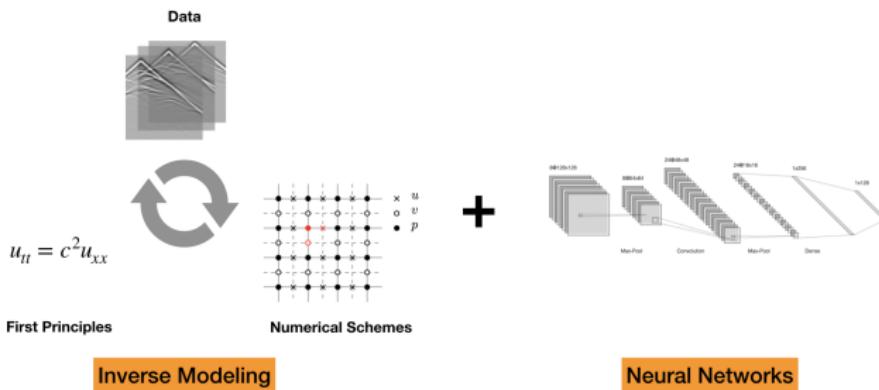
The candidate solution space is **infinite dimensional**.

Machine Learning for Computational Engineering

$$\min_{\theta} L_h(u_h) \quad \text{s.t. } F_h(\mathbf{NN}_{\theta}, u_h) = 0 \leftarrow \text{Solved numerically}$$

- ① Use a deep neural network to approximate the (high dimensional) unknown function;
- ② Solve u_h from the physical constraint using a numerical PDE solver;
- ③ Apply an unconstrained optimizer to the reduced problem

$$\min_{\theta} L_h(u_h(\theta))$$

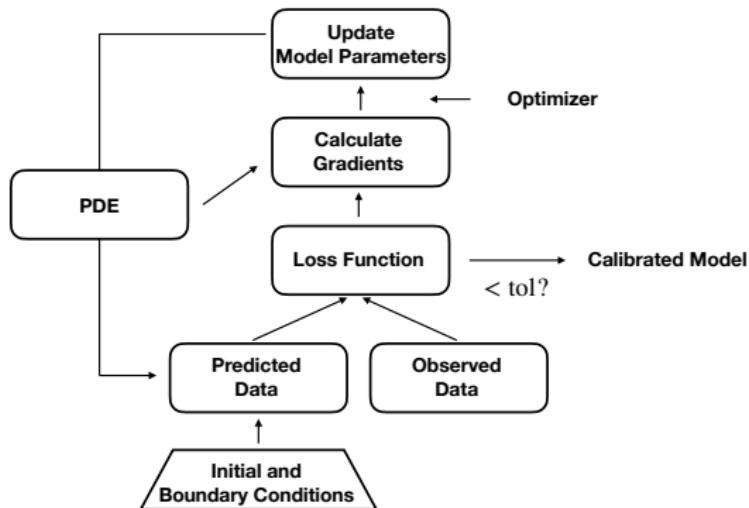


Gradient Based Optimization

$$\min_{\theta} L_h(u_h(\theta))$$

- Steepest descent method:

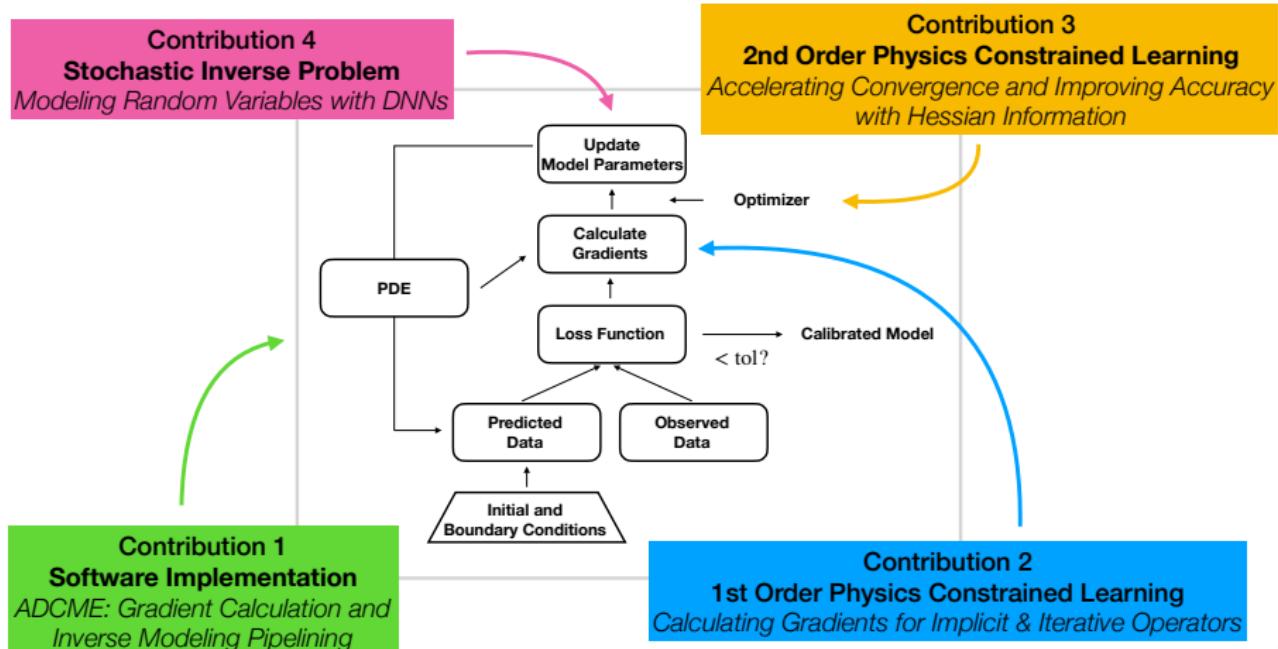
$$\theta_{k+1} \leftarrow \theta_k - \alpha_k \nabla_{\theta} L_h(u_h(\theta_k))$$

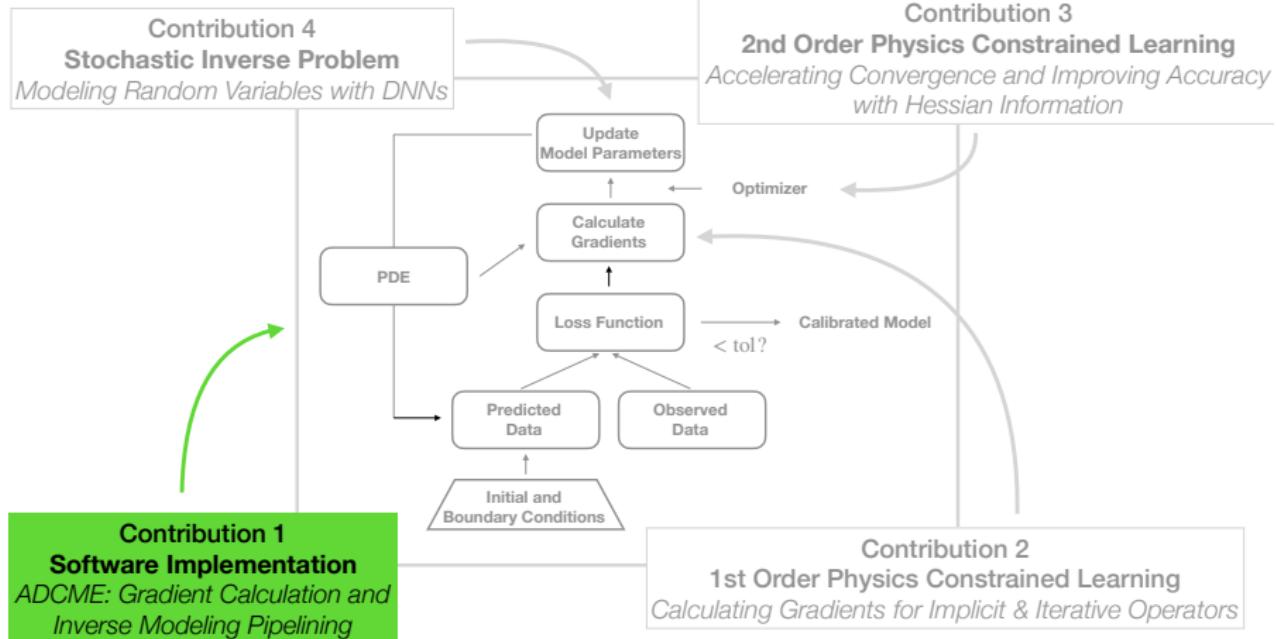


Contributions

Goal

Develop algorithms and tools for solving inverse problems by combining DNNs and numerical PDE solvers.





Ecosystem for Inverse Modeling

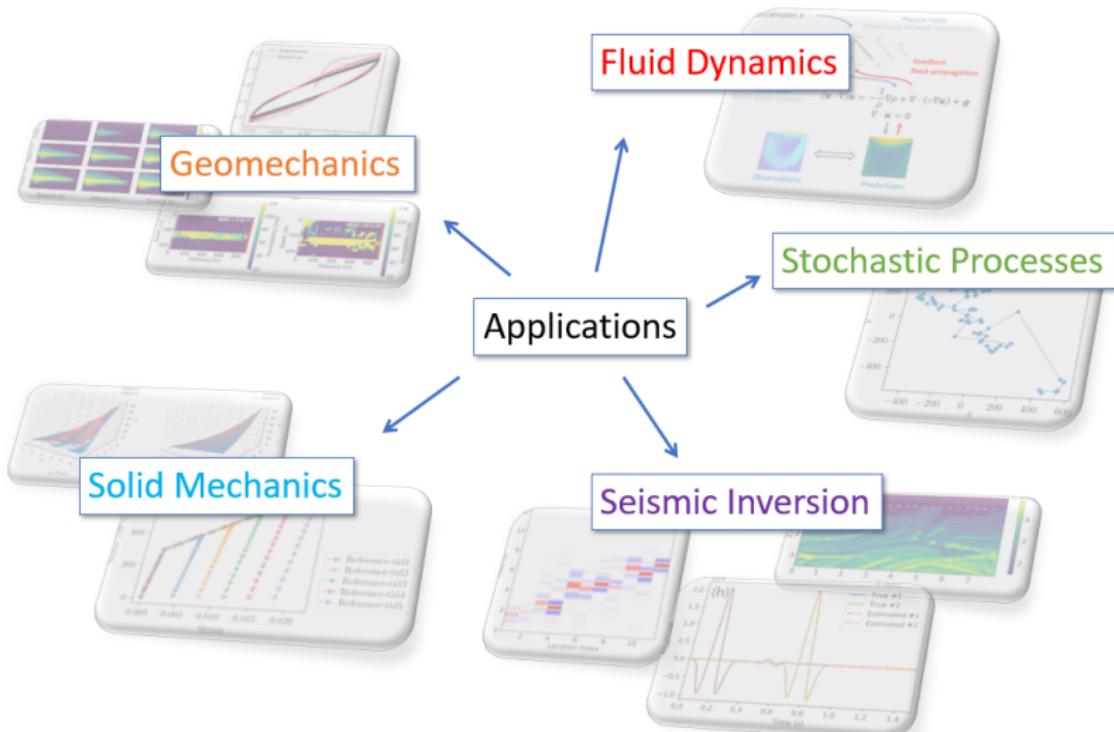
The figure displays a collage of software documentation pages, highlighting the interconnected nature of the ecosystem for inverse modeling:

- ADSeismic**: An open-source high-performance package for general seismic inversion problems. It includes a Getting Started guide and a detailed Documentation page.
- AdFem**: A software for finite element analysis, featuring a Documentation page and a GitHub repository.
- ADCME**: A library for automatic differentiation in scientific computing, with a Documentation page and a GitHub repository.
- SeisML**: A machine learning framework for seismic data processing, with a Documentation page and a GitHub repository.

The ADCME page provides an overview of its capabilities, including gradient-based optimization techniques, and links to various tutorials and examples. The SeisML page details its computational graph and forward-backward pass mechanisms. The AdFem page shows how it can be used for finite element simulations. The ADSeismic page integrates all these components to demonstrate how they work together for solving inverse modeling problems.

Documentation

Applications



See the publication list at: <https://github.com/kailaix/ADCME.jl>

Automatic Differentiation

Bridging the technical gap between deep learning and inverse modeling:

Mathematical Fact

Back-propagation

||

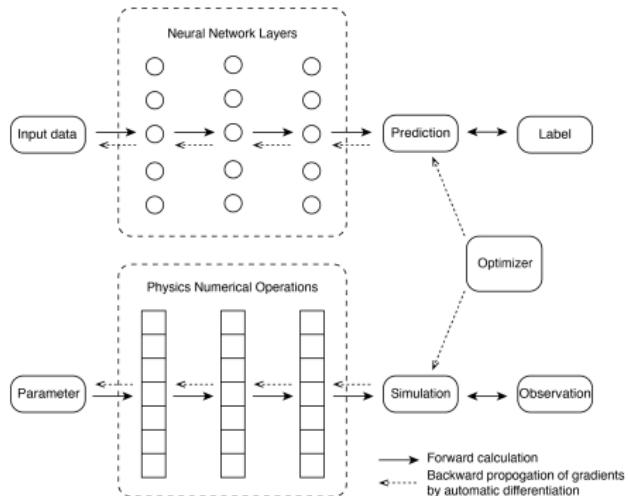
Reverse-mode

Automatic Differentiation

||

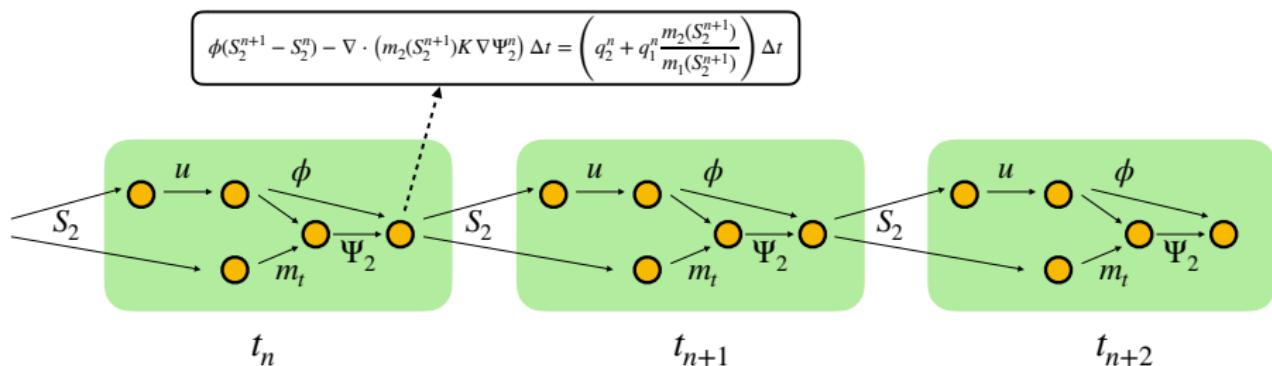
Discrete

Adjoint-State Method

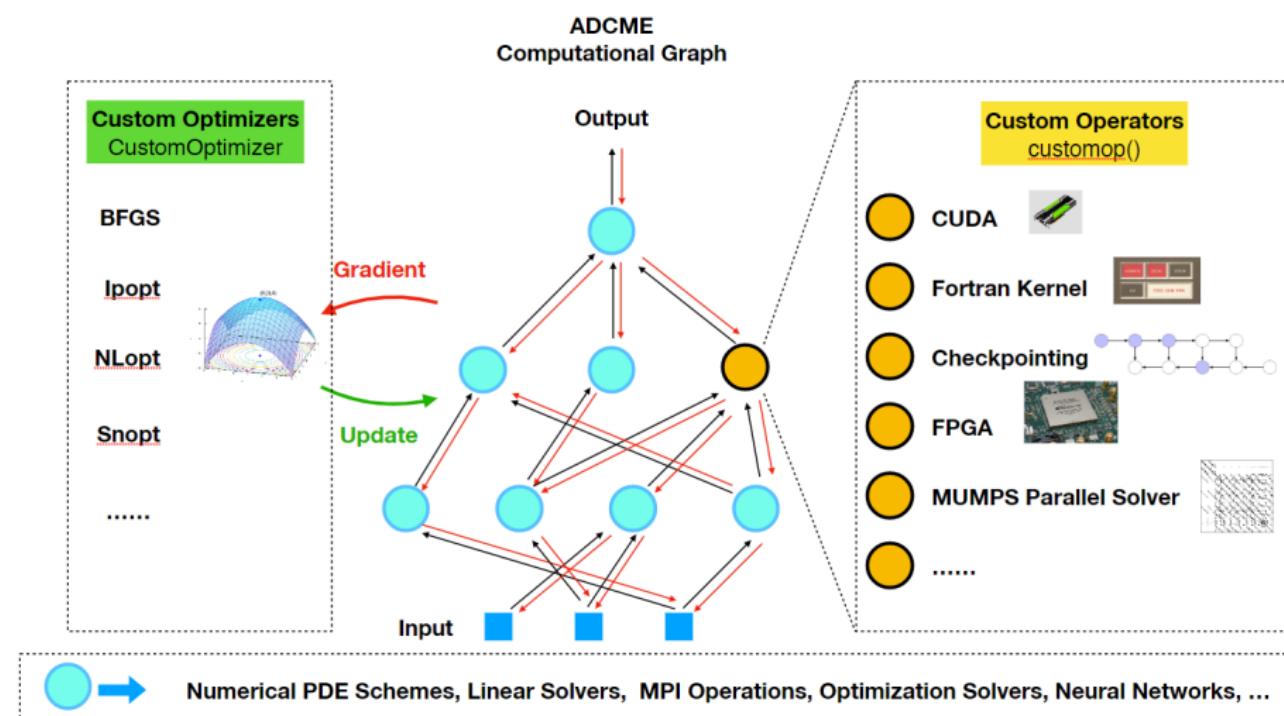


Computational Graph for Numerical Schemes

- To leverage automatic differentiation for inverse modeling, we need to express the numerical schemes in the “AD language”: computational graph.
- No matter how complicated a numerical scheme is, it can be decomposed into a collection of operators that are interlinked via state variable dependencies.



ADCME: Computational-Graph-based Numerical Simulation



How ADCME works

- ADCME translates your numerical simulation codes to computational graph and then the computations are delegated to a heterogeneous task-based parallel computing environment through TensorFlow runtime.

```
div  $\sigma(u) = f(x)$        $x \in \Omega$ 
 $\sigma(u) = C\varepsilon(u)$ 
 $u(x) = u_0(x)$        $x \in \Gamma_u$ 
 $\sigma(x)n(x) = t(x)$        $x \in \Gamma_n$ 

mesh = Mesh(50, 50, 1/50, degree=2)

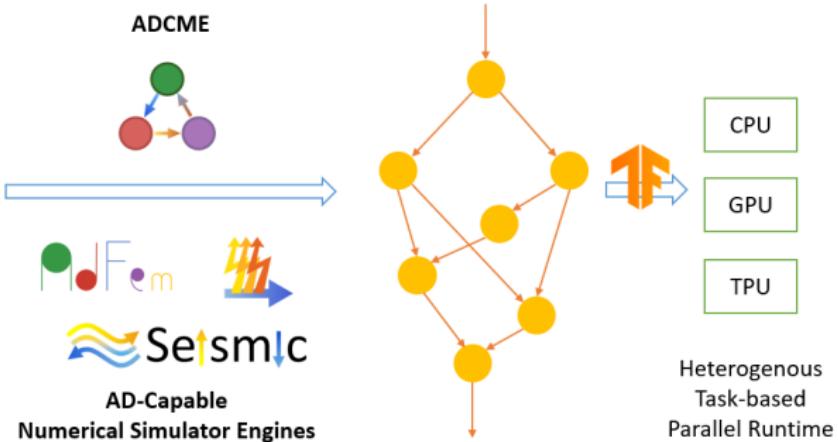
left = bcond((x,y)->x<1e-5, mesh)
right = bcond((x,y1,x2,y2)->(x2>0.049-1e-5) && (x2<0.050-1e-5), mesh)

t1 = eval_f_on_boundary_edge((x,y)->1.0e-4, right, mesh)
t2 = eval_f_on_boundary_edge((x,y)->0.0, right, mesh)
rhs = compute_fem_traction_term(t1, t2, right, mesh)

mu = 0.3
X = gauss_nodes(mesh)
E = abs(f(x, [20, 20, 20, 1]))>squeeze
# E = constant(eval_f_on_gauss_pts(f, mesh))

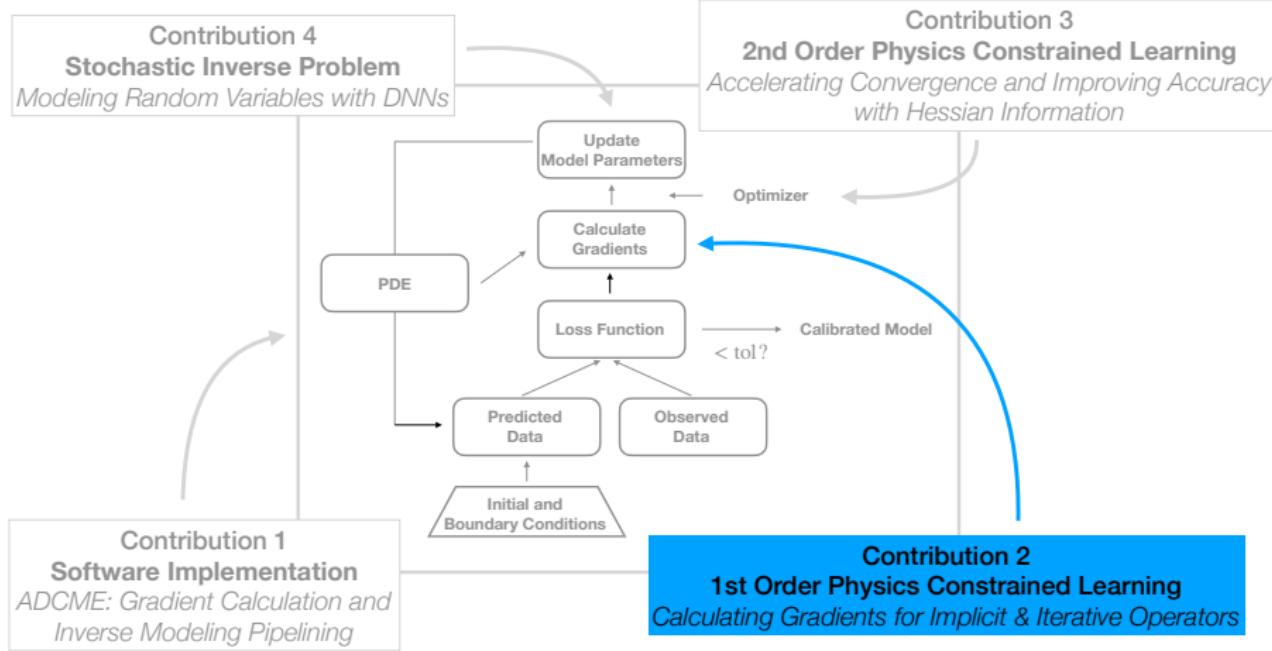
D = compute_plane_stress_matrix(X, mu*ones(gauss(mesh)))
K = compute_fem_stiffness_matrix(D, mesh)

bdval = [eval_f_on_boundary_node((x,y)->0.0, left, mesh);
         eval_f_on_boundary_node((x,y)->0.0, left, mesh)]
DOF = [left;left+mesh.ndof]
K, rhs = import_Dirichlet_boundary_conditions(K, rhs, DOF, bdval)
u = K\rhs
```



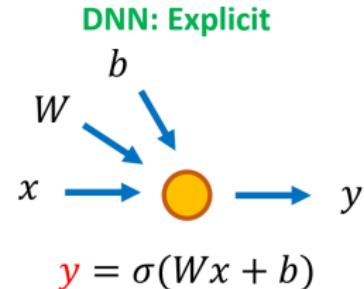
Summary

- Mathematically equivalent techniques for calculating gradients:
 - gradient back-propagation (DNN)
 - discrete adjoint-state methods (PDE)
 - reverse-mode automatic differentiation
- Computational graphs bridge the gap between gradient calculations in numerical PDE solvers and DNNs.
- ADCME extends the capability of TensorFlow to PDE solvers, providing users a single piece of software for numerical simulations, deep learning, and optimization.



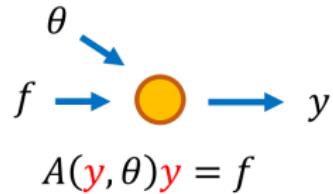
Motivation

- Most AD frameworks only deal with **explicit operators**, i.e., the functions that has analytical derivatives, or composition of these functions.
- Many scientific computing algorithms are **iterative** or **implicit** in nature.



Linear/Nonlinear	Explicit/Implicit	Expression
Linear	Explicit	$y = Ax$
Nonlinear	Explicit	$y = F(x)$
Linear	Implicit	$Ay = x$
Nonlinear	Implicit	$F(x, y) = 0$

Numerical Schemes:
Implicit, Iterative



Example

- Consider a function $f : x \rightarrow y$, which is implicitly defined by

$$F(x, y) = x^3 - (y^3 + y) = 0$$

If not using the cubic formula for finding the roots, the forward computation consists of iterative algorithms, such as the Newton's method and bisection method

```
y0 ← 0  
k ← 0  
while |F(x, yk)| > ε do  
    δk ← F(x, yk)/Fy'(x, yk)  
    yk+1 ← yk - δk  
    k ← k + 1  
end while  
Return yk
```

```
I ← -M, r ← M, m ← 0  
while |F(x, m)| > ε do  
    c ←  $\frac{a+b}{2}$   
    if F(x, m) > 0 then  
        a ← m  
    else  
        b ← m  
    end if  
end while  
Return c
```

Example

- An efficient way to do automatic differentiation is to apply the **implicit function theorem**. For our example, $F(x, y) = x^3 - (y^3 + y) = 0$; treat y as a function of x and take the derivative on both sides

$$3x^2 - 3y(x)^2y'(x) - y'(x) = 0 \Rightarrow y'(x) = \frac{3x^2}{3y^2 + 1}$$

The above gradient is **exact**.

Can we apply the same idea to inverse modeling?

Physics Constrained Learning (PCL)

$$\min_{\theta} L_h(u_h) \quad \text{s.t. } F_h(\theta, u_h) = 0$$

- Assume that we solve for $u_h = G_h(\theta)$ with $F_h(\theta, u_h) = 0$, and then

$$\tilde{L}_h(\theta) = L_h(G_h(\theta))$$

- Applying the **implicit function theorem**

$$\frac{\partial F_h(\theta, u_h)}{\partial \theta} + \frac{\partial F_h(\theta, u_h)}{\partial u_h} \frac{\partial G_h(\theta)}{\partial \theta} = 0 \Rightarrow \frac{\partial G_h(\theta)}{\partial \theta} = - \left(\frac{\partial F_h(\theta, u_h)}{\partial u_h} \right)^{-1} \frac{\partial F_h(\theta, u_h)}{\partial \theta}$$

- Finally we have

$$\boxed{\frac{\partial \tilde{L}_h(\theta)}{\partial \theta} = \frac{\partial L_h(u_h)}{\partial u_h} \frac{\partial G_h(\theta)}{\partial \theta} = - \frac{\partial L_h(u_h)}{\partial u_h} \left(\frac{\partial F_h(\theta, u_h)}{\partial u_h} \Big|_{u_h=G_h(\theta)} \right)^{-1} \frac{\partial F_h(\theta, u_h)}{\partial \theta} \Big|_{u_h=G_h(\theta)}}$$

Penalty Methods

$$\min_{\mathbf{f}} L_h(u_h) \quad \text{s.t. } F_h(\mathbf{f}, u_h) = 0$$

- **Penalty Method:** parametrize f with f_θ (DNNs, linear finite element basis, radial basis functions, etc.) and incorporate the physical constraint as a **penalty term** (regularization, prior, ...) in the loss function.

$$\min_{\theta, u_h} L_h(u_h) + \lambda \|F_h(f_\theta, u_h)\|_2^2$$

- + Easy to implement (no need for differentiating numerical solvers)
- May not satisfy physical constraint $F_h(f_\theta, u_h) = 0$ accurately;
- High dimensional optimization problem; both θ and u_h are variables.

Physics Constrained Learning for Stiff Problems

- PCL is superior for stiff problems.
- Consider a model problem

$$\min_{\theta} \|u - u_0\|_2^2 \quad \text{s.t. } Au = \theta y$$

$$\text{PCL : } \min_{\theta} \tilde{L}_h(\theta) = \|\theta A^{-1}y - u_0\|_2^2 = (\theta - 1)^2 \|u_0\|_2^2$$

$$\text{Penalty Method : } \min_{\theta, u_h} \tilde{L}_h(\theta, u_h) = \|u_h - u_0\|_2^2 + \lambda \|Au_h - \theta y\|_2^2$$

Theorem

The condition number of A_λ is

$$\liminf_{\lambda \rightarrow \infty} \kappa(A_\lambda) = \kappa(A)^2, \quad A_\lambda = \begin{bmatrix} I & 0 \\ \sqrt{\lambda}A & -\sqrt{\lambda}y \end{bmatrix}, \quad y = \begin{bmatrix} u_0 \\ 0 \end{bmatrix}$$

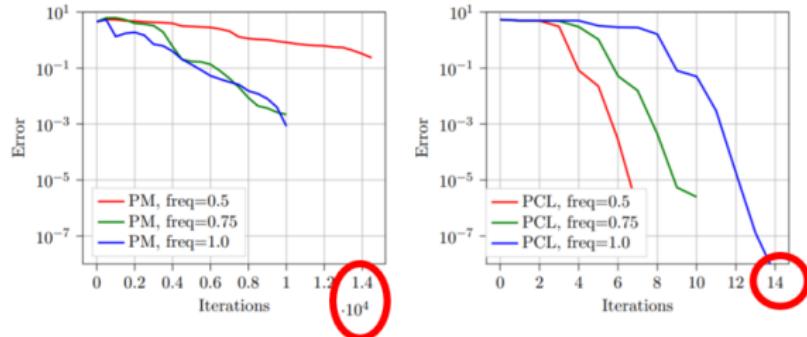
and therefore, the condition number of the unconstrained optimization problem from the penalty method is equal to the square of the condition number of the PCL asymptotically.

Physics Constrained Learning for Stiff Problems

Parameter Inverse Problem

$$\Delta u + k^2 g(x)u = 0$$
$$g(x) = 5x^2 + 2y^2$$

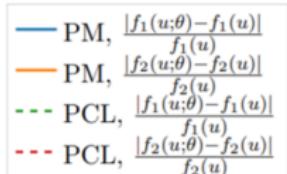
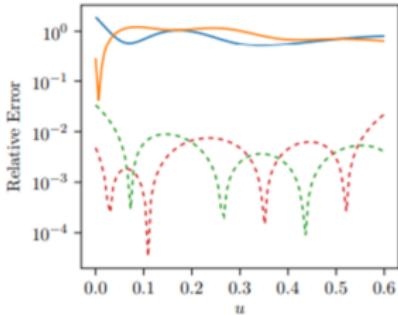
$$g_{\theta}(x) = \theta_1 x^2 + \theta_2 y^2 + \theta_3 xy \\ + \theta_4 x + \theta_5 y + \theta_6$$



Approximate Unknown Functions using DNNs

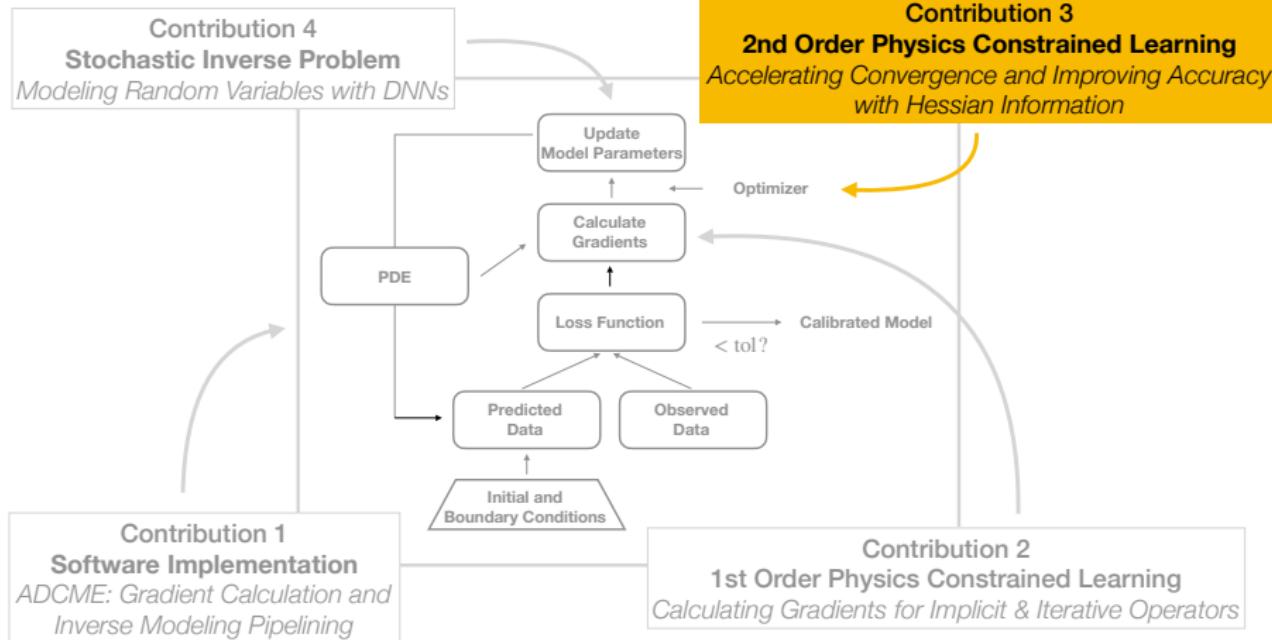
$$-\nabla \cdot (\mathbf{f}(\mathbf{u}) \nabla \mathbf{u}) = h(\mathbf{x})$$

$$\mathbf{f}(\mathbf{u}) = \begin{bmatrix} NN(u; \theta_1) & 0 \\ 0 & NN(u; \theta_2) \end{bmatrix}$$

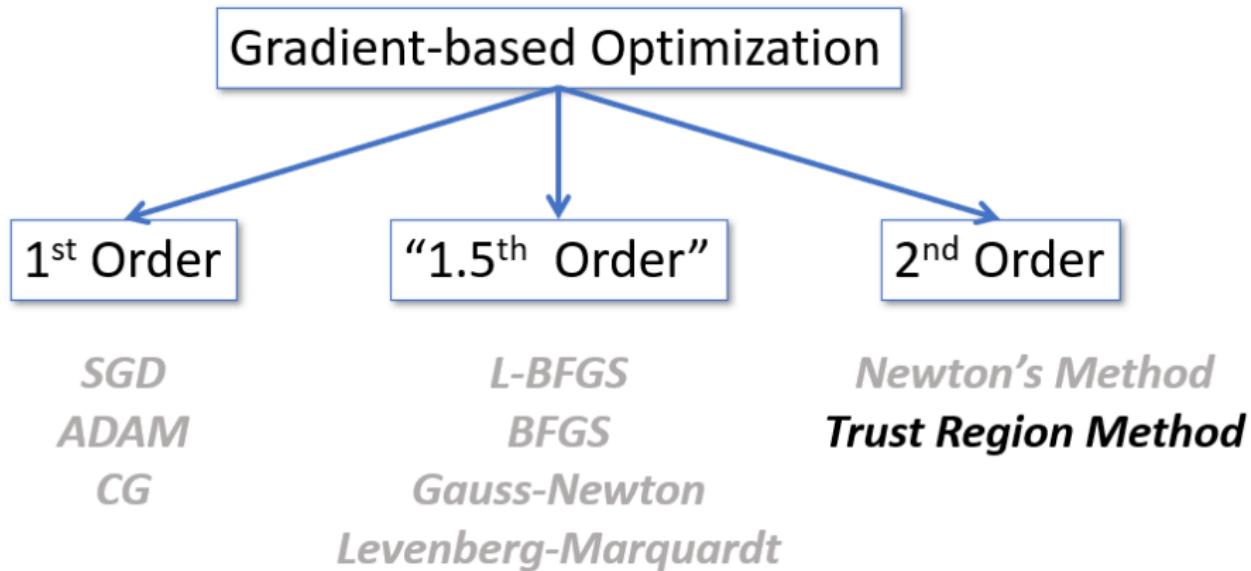


Summary

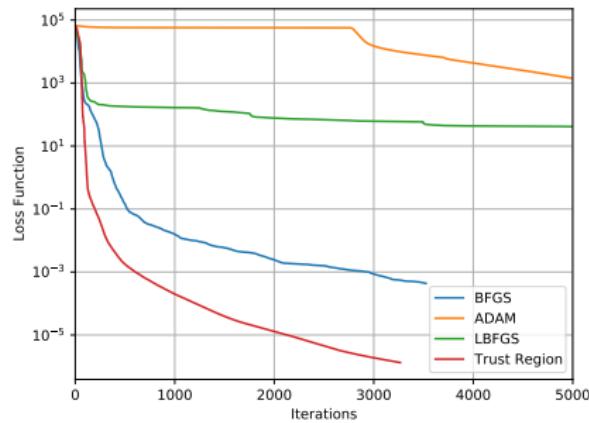
- Implicit and iterative operators are ubiquitous in numerical PDE solvers. These operators are insufficiently treated in deep learning software and frameworks.
- PCL helps you calculate gradients of implicit/iterative operators efficiently.
- PCL leads to faster convergence and better accuracy compared to penalty methods for stiff problems.



Motivation



Overview



Goal

Accelerate convergence and improve accuracy with Hessian information

Challenge

Calculate Hessians for coupled systems of PDEs and DNNs

Trust Region vs. Line Search

Trust Region

- Approximate $f(x_k + p)$ by a model quadratic function

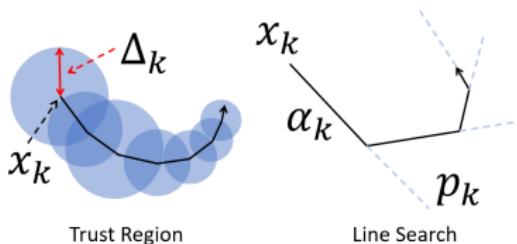
$$m_k(p) = f_k + g_k^T p + \frac{1}{2} p^T B_k p$$

$$f_k = f(x_k), g_k = \nabla f(x_k), B_k = \nabla^2 f(x_k)$$

- Solve the optimization problem within a trust region $\|p\| \leq \Delta_k$

$$p_k = \arg \min_p m_k(p) \quad \text{s.t. } \|p\| \leq \Delta_k$$

- If decrease in $f(x_k + p_k)$ is sufficient, then update the state $x_{k+1} = x_k + p_k$; otherwise, $x_{k+1} = x_k$ and improve Δ_k .



Line Search

- Determine a descent direction p_k
- Determine a step size α_k that sufficiently reduces $f(x_k + \alpha_k p_k)$
- Update the state
$$x_{k+1} = x_k + \alpha_k p_k$$

Second-order Physics Constrained Learning

- Consider a composite function with a vector input x and scalar output

$$v = f(G(x)) \quad (1)$$

- Define

$$\begin{aligned}f_{,k}(y) &= \frac{\partial f(y)}{\partial y_k}, & f_{,kl}(y) &= \frac{\partial^2 f(y)}{\partial y_k \partial y_l} \\G_{k,I}(x) &= \frac{\partial G_k(x)}{\partial x_I}, & G_{k,Ir}(x) &= \frac{\partial^2 G_k(x)}{\partial x_I \partial x_r}\end{aligned}$$

- Differentiate Equation (1) with respect to x_i

$$\frac{\partial v}{\partial x_i} = f_{,k} G_{k,i} \quad (2)$$

- Differentiate Equation (2) with respect to x_j

$$\frac{\partial^2 v}{\partial x_i \partial x_j} = f_{,kr} G_{k,i} G_{r,j} + f_{,k} G_{k,ij}$$

Second-order Physics Constrained Learning

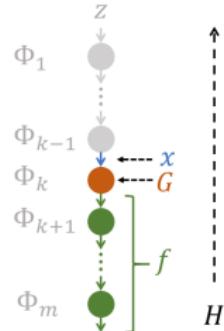
In the vector form,

$$\nabla^2 v = (\nabla G)^T \nabla^2 f(\nabla G) + \nabla^2(\bar{G}^T G) \quad \bar{G} = \nabla f$$

- Consider a function composed of a sequence of computations

$$v = \Phi_m(\Phi_{m-1}(\cdots(\Phi_1(z))))$$

- 1: Initialize $H \leftarrow 0$
- 2: **for** $k = m - 1, m - 2, \dots, 1$ **do**
- 3: Define $f := \Phi_m(\Phi_{m-1}(\cdots(\Phi_{k+1}(\cdot))))$, $G := \Phi_k$
- 4: Calculate the gradient (Jacobian) $J \leftarrow \nabla G$
- 5: Extract \bar{G} from the saved gradient back-propagation data.
- 6: Calculate $Z = \nabla^2(\bar{G}^T G)$
- 7: Update $H \leftarrow J^T H J + Z$
- 8: **end for**



Numerical Benchmark

- We consider the heat equation in $\Omega = [0, 1]^2$

$$\frac{\partial u}{\partial t} = \nabla \cdot (\kappa(x, y) \nabla u) + f(x, y) \quad x \in \Omega$$

$$u(x, y, 0) = x(1 - x)y^2(1 - y)^2 \quad (x, y) \in \Omega$$

$$u(x, y, t) = 0 \quad (x, y) \in \partial\Omega$$

- The diffusivity coefficient κ and exact solution u are given by

$$\kappa(x, y) = 2x^2 - 1.05x^4 + x^6 + xy + y^2$$

$$u(x, y, t) = x(1 - x)y^2(1 - y)^2 e^{-t}$$

- We learn a DNN approximation to κ using full-field observations of u

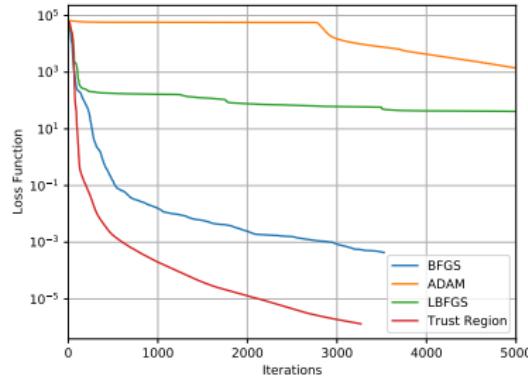
$$\kappa(x, y) \approx \text{NN}_\theta(x, y)$$

Convergence

- The optimization problem is given by

$$\min_{\theta} L(\theta) = \sum_n \sum_{i,j} \left(\frac{u_{i,j}^{n+1} - u_{i,j}^n}{\Delta t} - F_{i,j}(u^{n+1}; \theta) - f_{i,j}^{n+1} \right)^2$$

Here $F_{i,j}(u^{n+1}; \theta)$ is the 4-point finite difference approximation to the Laplacian $\nabla \cdot (\text{NN}_{\theta} \nabla u)$.



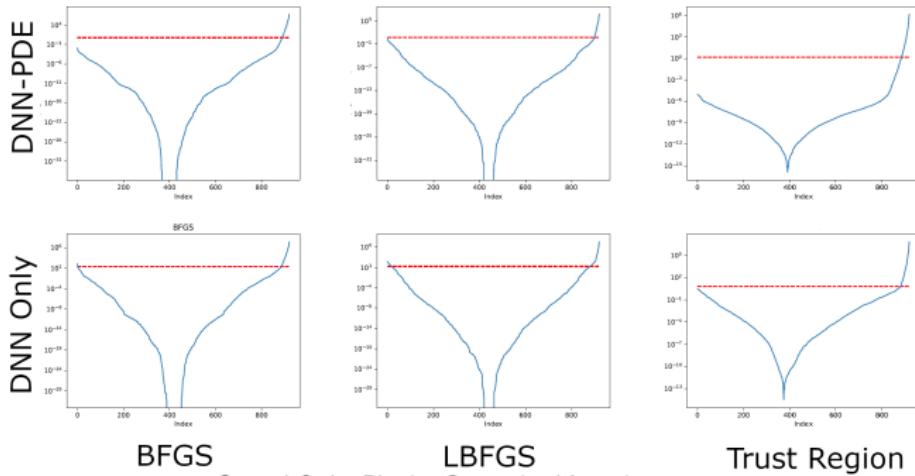
Effect of PDEs

NN_θ → (PDE Solver) → Loss Function

- Consider the loss function excluding the effects of PDEs

$$I(\theta) = \sum_{i,j} (\text{NN}_\theta(x_{i,j}, y_{i,j}) - \kappa(x_{i,j}, y_{i,j}))^2$$

- Eigenvalue magnitudes of $\nabla^2 L(\theta)$ and $\nabla^2 I(\theta)$



Effect of PDEs

- Most of the eigenvalue directions at the local landscape of loss functions are “flat” \Rightarrow “effective degrees of freedom (DOFs)”.
- Physical constraints (PDEs) further cannibalize effective DOFs:

	BFGS	LBFGS	Trust Region
DNN-PDE	31	22	35
DNN Only	34	41	38

Effect of Widths and Depths

- The ratio of zero eigenvalues **increases** as
 - the number of hidden layers increase for a fixed number (20) of neurons per layer

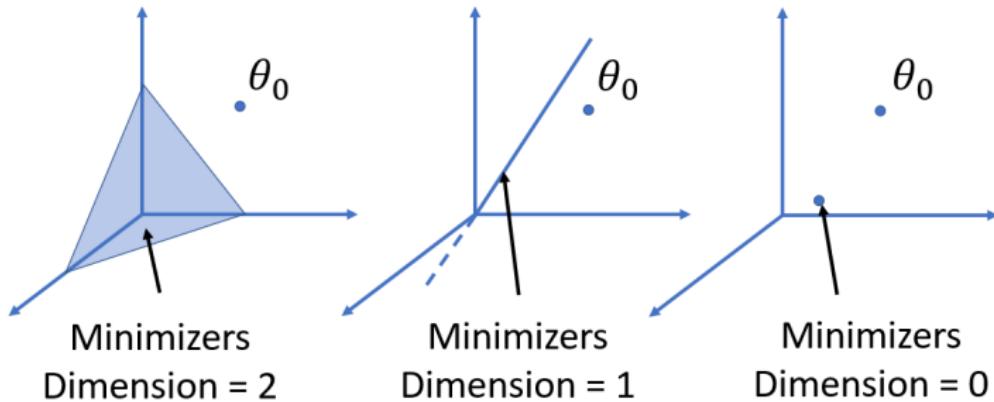
# Hidden Layers	LBFGS	BFGS	Trust Region
1	76.54	72.84	77.78
2	98.2	94.41	93.21
3	98.7	98.15	96.09

- the number of neurons per layer increases for a fixed number (3) of hidden layers

# Neurons per Layer	LBFGS	BFGS	Trust Region
5	93.83	85.19	69.14
10	97.7	83.52	89.66
20	96.2	97.39	96.42

Effect of Widths and Depths: conjecture

- Implications for overparametrization: **the minimizer lies on a relatively higher dimensional manifold of the parameter space.**
- Conjecture: overparameterization makes the optimization easier due to a larger chance of hitting the minimizer manifold.



Summary

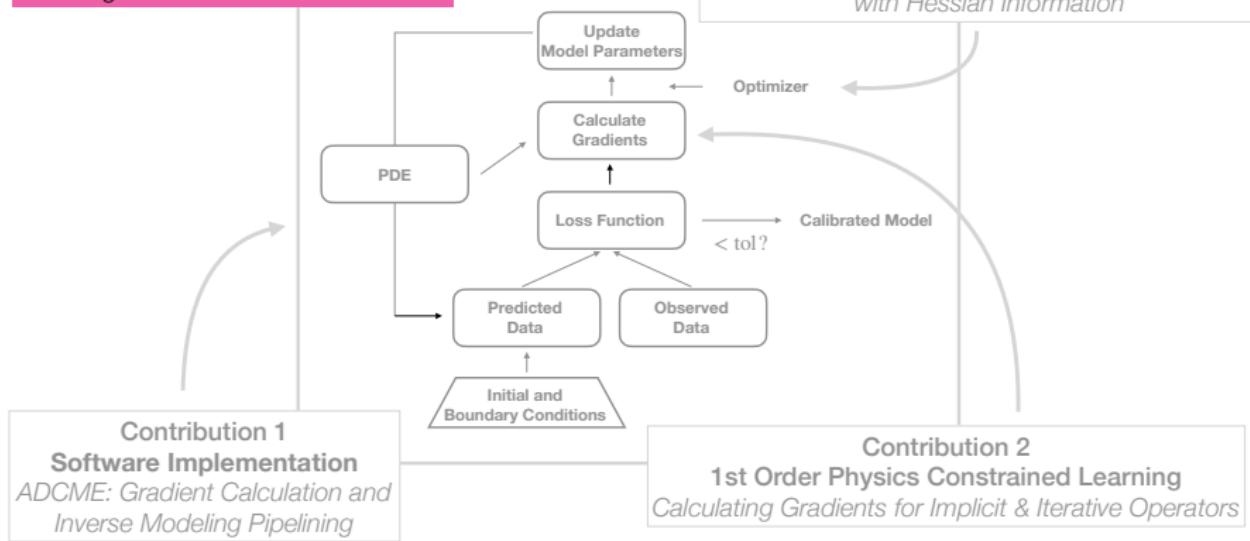
- Trust region methods converge significantly faster compared to first order/quasi second order methods by leveraging Hessian information.
- Second order physics constrained learning helps you calculate Hessian matrices efficiently.
- The local minimum of DNNs have small effective degrees of freedom compared to DNN sizes.

Contribution 4 Stochastic Inverse Problem

Modeling Random Variables with DNNs

Contribution 3 2nd Order Physics Constrained Learning

Accelerating Convergence and Improving Accuracy with Hessian Information

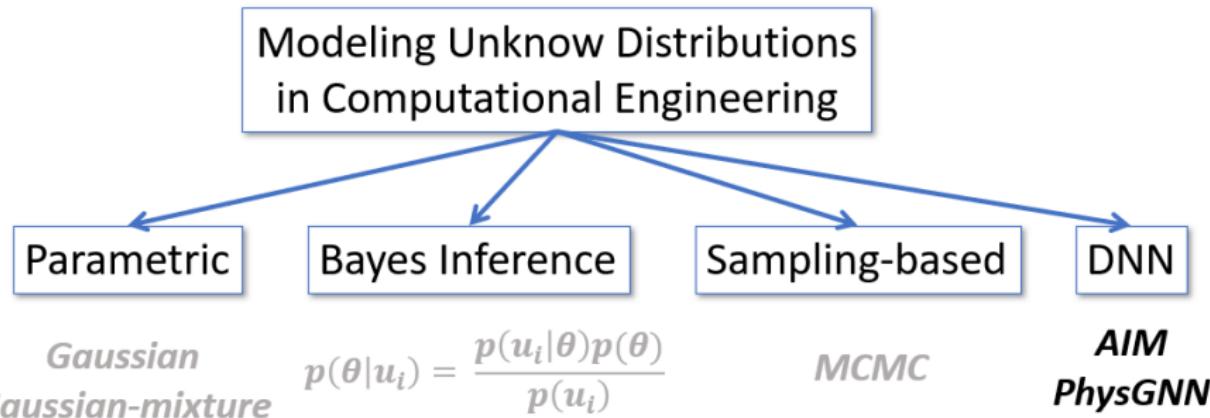


Motivation

$$\min_{\mathbf{f}} L_h(u_h) \quad \text{s.t. } F_h(\mathbf{f}, u_h) = 0$$

- What if f is a random variable?
 - uncertainty quantification
- What if F_h is a stochastic model?
 - stochastic differential equations (SDE)

Methodologies



Modeling Stochasticity using Deep Neural Networks

- We consider a map F

$$F : (w, \theta) \mapsto u(\cdot, \cdot)$$

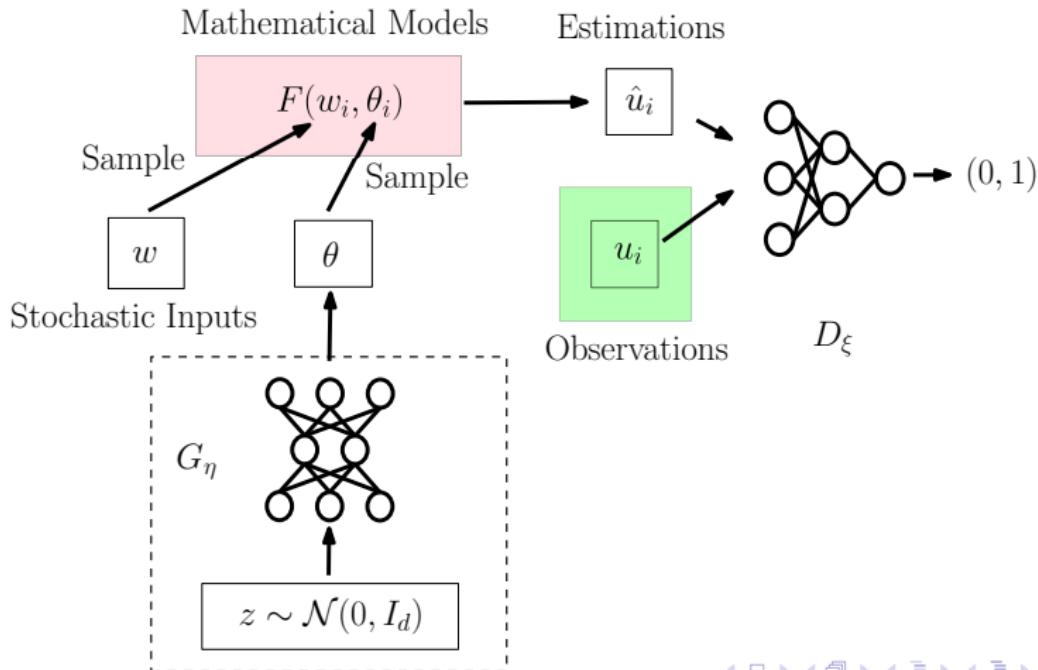
- w is sampled from a **know** stochastic process (stochasticity intrinsic to the model).
- θ is a random variable of interest.
- $u(x, t)$ is the output of the mapping, which depends on the location x and time t . u is also a random variable.
- Examples:
 - CIR process (a short-term interest rate model)

$$dr_t = \kappa(\tau - r_t)dt + \sqrt{r_t}\sigma dW_t$$

Here $\theta = (\kappa, \tau, \sigma)$, $w = W_t$, and $u = r_t$.

Adversarial Inverse Modeling (AIM)

- Approximate the unknown distribution θ with a DNN G_η (generator);
- A discriminator D_ξ tries to distinguish between generated and observed samples.



Adversarial Inverse Modeling

- Inspired by generative adversarial nets (GAN): alternatively minimizing discriminator and generator loss functions
 - Kullback-Leibler (KL) GAN

$$L^G(\{\hat{u}_i\}; \eta) = \frac{1}{n} \sum_{i=1}^n \log \frac{1 - D_\xi(\hat{u}_i(\eta))}{D_\xi(\hat{u}_i(\eta))}$$

$$L^D(\{u_i\}, \{\hat{u}_i\}; \xi) = - \frac{1}{n} \sum_{i=1}^n (\log D_\xi(\hat{u}_i) + \log(1 - D_\xi(u_i)))$$

- Wasserstein GAN

$$L^G(\{\hat{u}_i\}; \eta) = \frac{1}{n} \sum_{i=1}^n D_\xi(\hat{u}_i(\eta))$$

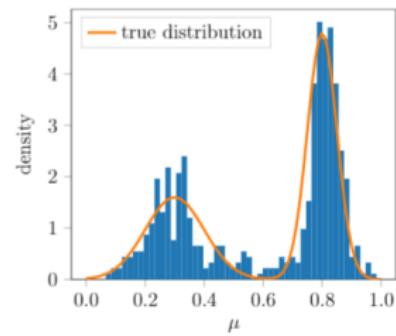
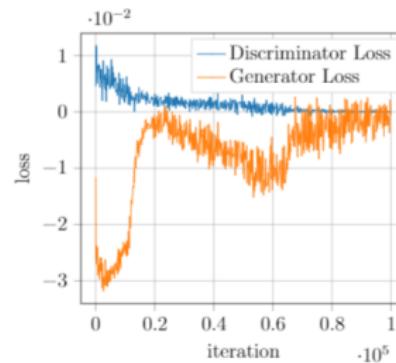
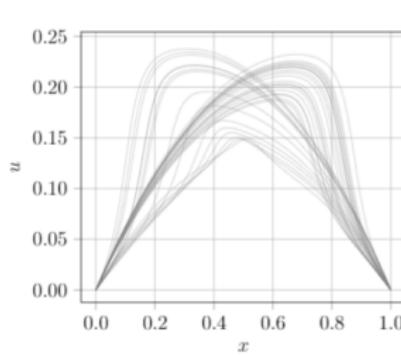
$$L^D(\{u_i\}, \{\hat{u}_i\}; \xi) = - \frac{1}{n} \sum_{i=1}^n D_\xi(u_i)$$

Numerical Benchmarks

- We consider a Poisson's equation

$$\begin{cases} -\nabla \cdot (a(x)\nabla u(x)) = 1 & x \in (0, 1) \\ u(0) = u(1) = 0 & \text{otherwise} \end{cases}$$

$$a(x) = 1 - 0.9 \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

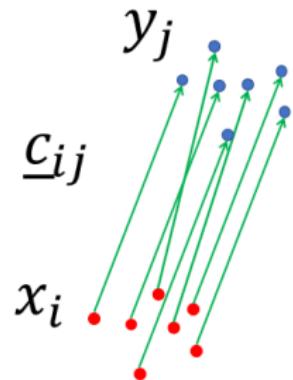


Optimal Transport

- Solve the transport matrix from a linear programming problem

$$\min_{p_{ij}} W(p; x, y) = \sum_{i=1}^n \sum_{j=1}^m p_{ij} c_{ij}$$

$$\text{s.t. } \sum_{j=1}^m p_{ij} = \frac{1}{n}, \quad \sum_{i=1}^n p_{ij} = \frac{1}{m}, \quad 0 \leq p_{ij} \leq 1$$



- Discrete Wasserstein distance between two sets of samples:

$$D(x, y) = W(p^*; x, y) = \sum_{1 \leq i \leq n, 1 \leq j \leq m} p_{ij}^* \|x_i - y_j\|_1$$

AD for Discrete Wasserstein Distance

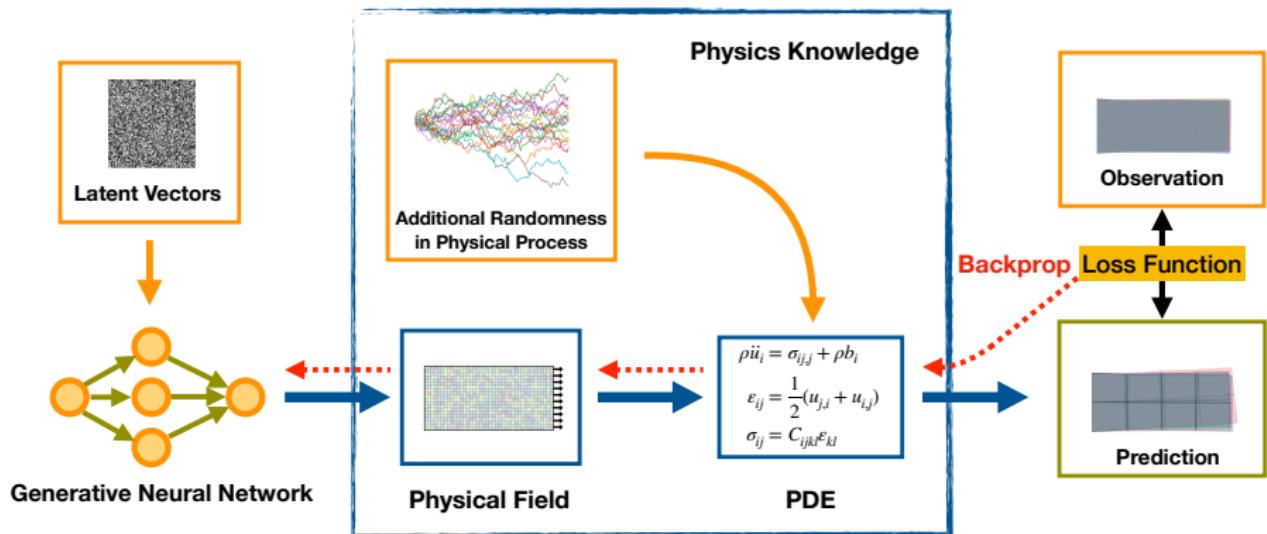
Theorem

Assume that $\mathbf{p}^*([\underline{c}_{ij}])$ is the optimal solution to the linear programming problem given the cost matrix $[\underline{c}_{ij}]$, then we have

$$\frac{\partial W(\mathbf{p}^*([\underline{c}_{ij}]))}{\partial \underline{c}_{ij}} = p_{ij}^*$$

- We have a “**handpicked discriminator**”.
- Numerical implementation is very similar to OptNet (Amos, Brandon, and J. Zico Kolter, PMLR 2017):
 - The forward computation constitutes solving an optimization problem.

Physics Generative Neural Network (PhysGNN)



Numerical Benchmarks

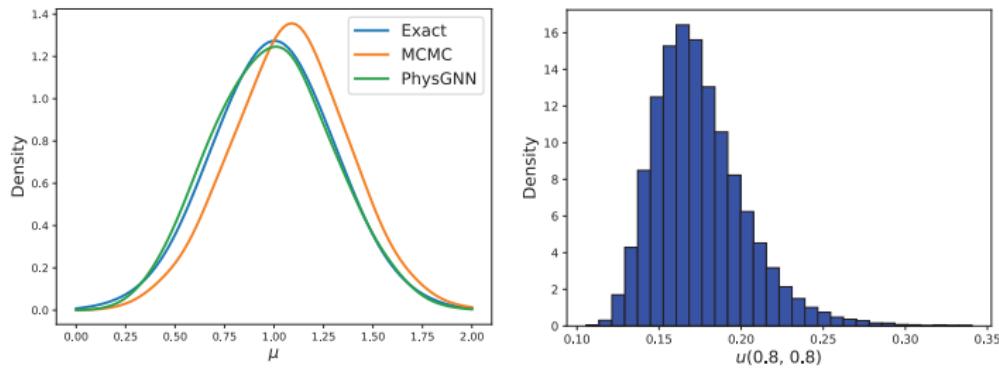
- We consider a 2D Poisson's equation

$$\nabla \cdot (\kappa \nabla u) = 2\pi^2 \sin(\pi x) \sin(\pi y) \quad (x, y) \in (0, 1)^2$$

$$u(x, y) = 0 \quad (x, y) \in \partial(0, 1)^2$$

$$\kappa = |v|, \quad v \sim \mathcal{N}(1.0, 0.3^2)$$

- Observation: $u(0.8, 0.8)$
- $\sim 20,000$ PDE solves for both MCMC and PhysGNN:



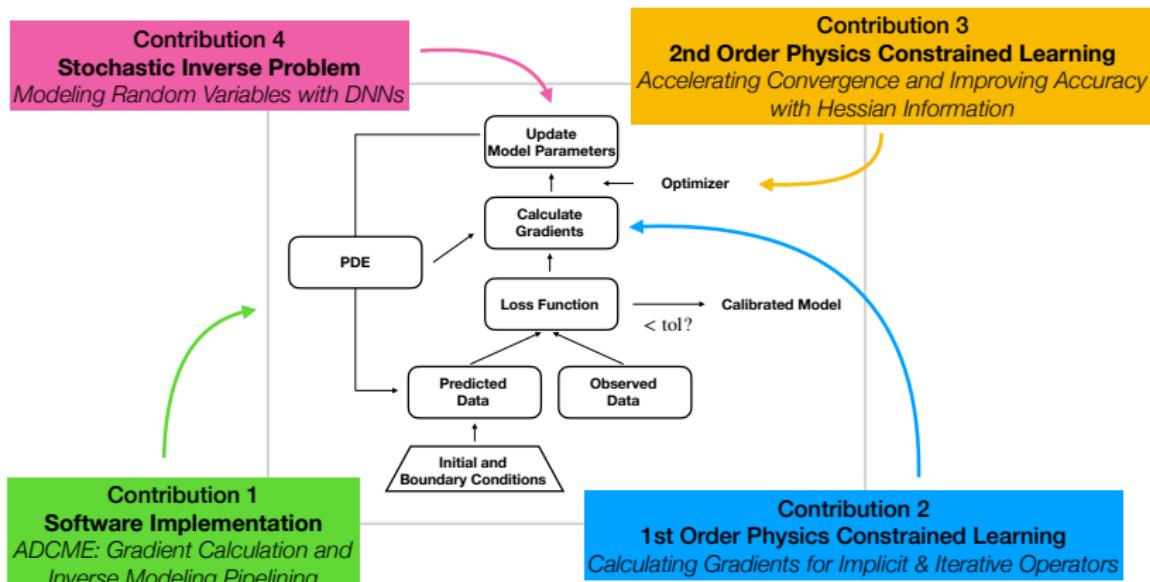
Summary

- Generative neural networks are used to approximate unknown distributions in a stochastic inverse problem.
- Training generative neural networks
 - Adversarial inverse modeling (AIM): solving a min-max optimization problem involving generator and discriminator nets
 - Physics Generative Neural Network (PhysGNN): minimizing a discrete Wasserstein distance

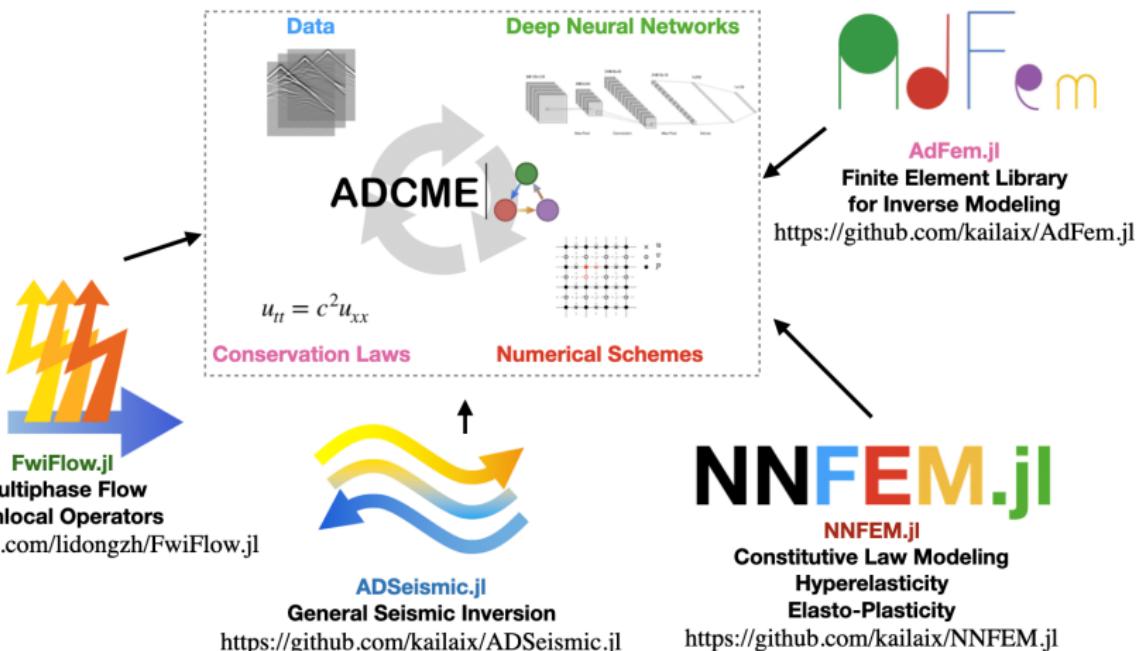
Conclusion

$$\min_f L_h(u_h) \quad \text{s.t. } F_h(f, u_h) = 0$$

- ✓ Develop algorithms and tools for solving inverse problems by combining DNNs and numerical PDE solvers.



A General Approach to Inverse Modeling



SURPRISE!